

Highlight

Note

## Publishing a Batch Inferencing Service

Just like any pipeline, you can publish a batch inferencing pipeline as a REST service:

```
published_pipeline = pipeline_run.publish_pipeline(name='Batch_Prediction_Pipeline',
                                                    description='Batch pipeline',
                                                    version='1.0')

rest_endpoint = published_pipeline.endpoint
```

Once published, you can use the service endpoint to initiate a batch inferencing job:

```
import requests

response = requests.post(rest_endpoint,
                        headers=auth_header,
                        json={"ExperimentName": "Batch_Prediction"})
run_id = response.json()["Id"]
```

You can also schedule the published pipeline to have it run automatically:

```
from azureml.pipeline.core import ScheduleRecurrence, Schedule

weekly = ScheduleRecurrence(frequency='Week', interval=1)
pipeline_schedule = Schedule.create(ws, name='Weekly Predictions',
                                    description='batch inferencing',
                                    pipeline_id=published_pipeline.id,
                                    experiment_name='Batch_Prediction',
                                    recurrence=weekly)
```

This document belongs to Wolfgang Kiesenhofer.  
wolfgang.kiesenhofer@outlook.com  
No unauthorized copies allowed!

This document belongs to  
wolfgang.kiesenhofer@outlook.com  
No unauthorized copies allowed!