

Highlight

Note

Azure Machine Learning and Azure Pipelines

Azure Pipelines is a commonly used mechanism for automating software build and release processes.

At its simplest, Azure Pipelines provides a way to define a process that runs on-demand or in response to a source control event (such as pushing an update to a code file). You can use Azure CLI commands or Python scripts in a pipeline, making it an ideal way to automate Azure Machine Learning operations; such as training, registering, or deploying a model.

Additionally, you can install the Machine Learning extension for Azure pipelines to provide explicit integration with Azure Machine Learning. This enables you to monitor a workspace for new model registration events, and run a pipeline in response – for example to check the predictive performance of the newly registered model, and if it performs better than a previously deployed model, deploy the new one. You can also leverage Azure Machine Learning specific tasks in an Azure pipeline to run a published Azure Machine Learning pipeline, profile a model (to determine the appropriate deployment environment based on expected utilization), or deploy a model as a service.

This document belongs to Wolfgang Kiesenhofer.
wolfgang.kiesenhofer@outlook.com
No unauthorized copies allowed!

This document belongs to Wolfgang Kiesenhofer.
wolfgang.kiesenhofer@outlook.com
No unauthorized copies allowed!

This document belongs to Wolfgang Kiesenhofer.
wolfgang.kiesenhofer@outlook.com
No unauthorized copies allowed!